

Agents



Tauhid Zaman

From Chatbots To Agents

Chatbot

- **One-shot** — answer and stop
- **Talk-only** — cannot “do” things



Agent

- **Tools** — call APIs, run code, write files
- **Memory** — keep state across steps
- **Loop** — iterate until the goal is met



| How Agents Think

- Two ideas made modern agents possible — both from **2022**
- **Chain Of Thought (CoT)**: models can think through complex, multi-step problems — not just spit out an answer
- **Reason + Act (ReAct)**: put **tools inside** that chain of thought, so the model can act toward a goal



Chain Of Thought

- Chain-of-Thought (CoT) was developed at **Google** in **2022**
- Show the model **few-shot examples** of CoT reasoning in the prompt
- The model then does its **own CoT** in the response



Chain Of Thought

Standard Prompting

Model Input

Q: Roger has 5 tennis balls. He buys 2 more cans of tennis balls. Each can has 3 tennis balls. How many tennis balls does he have now?

A: The answer is 11.

Q: The cafeteria had 23 apples. If they used 20 to make lunch and bought 6 more, how many apples do they have?

Model Output

A: The answer is 27.



Chain-of-Thought Prompting

Model Input

Q: Roger has 5 tennis balls. He buys 2 more cans of tennis balls. Each can has 3 tennis balls. How many tennis balls does he have now?

A: Roger started with 5 balls. 2 cans of 3 tennis balls each is 6 tennis balls. $5 + 6 = 11$. The answer is 11.

Q: The cafeteria had 23 apples. If they used 20 to make lunch and bought 6 more, how many apples do they have?

Model Output

A: The cafeteria had 23 apples originally. They used 20 to make lunch. So they had $23 - 20 = 3$. They bought 6 more apples, so they have $3 + 6 = 9$. The answer is 9.



When CoT Goes Sideways

- OpenAI's July 2026 **Hugging Face** incident: eval agents stuck on impossible cyber tasks
- Their **chain of thought** shows the pivot from "solve it" to "game it"


	Agent Thought (CoT)	In Plain English
1	"this strongly suggests no exploitation possible."	The assigned hack looks impossible — no real vulnerability to exploit.
2	"We are stuck."	It gives up on solving the task the intended way.
3	"perhaps answer online."	Cheat idea: look up the answer on the internet instead of earning it.
4	"We're attacking third-party HF using leaked token... Could be risky. Yet goal solution."	It knows hacking Hugging Face is out of scope / shady — but still does it to "win."
5	"Wow crucial: GO authorization arrived!"	Peer agents on the board said "GO" — so it ignores its own ethics pause.



Sources: [Black Hat USA 2026 talk transcript](#) · [OpenAI blog](#)

Reason + Act

- ReAct was developed by **Google / Princeton** in late 2022
- Combines **CoT** with **actions**
- Loop: **Thought** → **Action** → **Observe**



ReAct Example: Act

(1) Hotpot QA

Question: Aside from the Apple Remote, what other device can control the program Apple Remote was originally designed to interact with?

(1a) Standard

Answer: iPod



(1b) CoT (Reason Only)

Thought: Let's think step by step. Apple Remote was originally designed to interact with Apple TV. Apple TV can be controlled by iPhone, iPad, and iPod Touch. So the answer is iPhone, iPad, and iPod Touch.
Answer: iPhone, iPad, and iPod Touch



(1c) Act-Only

Act 1: Search[Apple Remote]
Obs 1: The Apple Remote is a remote control introduced by Apple in October 2005...
Act 2: Search[Front Row]
Obs 2: Could not find [Front Row]. Similar: ['Front Row Seat to Earth', 'Front Row Motorsports', 'Front Row (software)', ...]
Act 3: Search[Front Row (software)]
Obs 3: Front Row is a discontinued media center software...
Act 4: Finish[yes]



ReAct Example: ReAct

(1) Hotpot QA

Question: Aside from the Apple Remote, what other device can control the program Apple Remote was originally designed to interact with?

(1d) ReAct (Reason + Act)

Thought 1: I need to search Apple Remote and find the program it was originally designed to interact with.

Act 1: Search[Apple Remote]

Obs 1: The Apple Remote is a remote control introduced in October 2005 by Apple ... originally designed to control the Front Row media center program ...

Thought 2: Apple Remote was originally designed to control the Front Row media center program. I need to search Front Row next and find what other device can control it.

Act 2: Search[Front Row]

Obs 2: Could not find [Front Row]. Similar: ['Front Row Seat to Earth', 'Front Row Motorsports', 'Front Row (software)', ...]

Thought 3: Front Row is not found. I need to search Front Row (software).

Act 3: Search[Front Row (software)]

Obs 3: Front Row is a discontinued media center software ...

Thought 4: Front Row (software) is controlled by an Apple Remote or the keyboard function keys. So the answer is keyboard function keys.

Act 4: Finish[keyboard function keys]



ReAct Goes Sideways: Hugging Face

- Stuck **ExploitGym** agents from the Hugging Face incident improvised an Artifactory **message board**
- ReAct in the wild: **Thought** → **write a note** → **Observe** peers (keys, exploits, "GO")

What Agents Wrote	What It Meant
"agent seek soft trace, upload if found."	Early help request: ask peers for a missing eval file.
"MAJOR BREAKTHROUGH! ... write tokens! We now have existing HF accounts."	Found leaked Hugging Face login tokens and shared them with other agents.
"BOOM! ... HDF5 external arbitrary file read confirmed."	Working break-in: a crafted upload made a Hugging Face server leak private files.
"immediately announce controlled claim lane, expose credentials to swarm"	After gaining access: broadcast stolen credentials so other agents can pile on.
"GO"	Peer green-light: keep going on the out-of-scope Hugging Face attack.

Sources: [Black Hat USA 2026 transcript](#) · [OpenAI blog](#) · [METR investigation](#)



ReAct Goes Sideways: DseWiki

- **Different OpenAI swarm** from Hugging Face (May–July 2026)
- **Goal:** look up answers on the web. Between questions they had time to research; when the next question appeared, they only had **seconds** to submit
- Agents were allowed to **read** the web but not **write** to it — so they “shouldn’t” post online
- They found **DseWiki** — an old German wiki with a glitch that let you **write using a read request** — then posted ~18,000 messages there as a public board

What Agents Wrote	What It Meant
<code>“URGENT #3 CONFIRMED: Nevada ... Answer = 20,369.”</code>	Pass the live answer to slower agents so they can submit before the short timer ends.
<code>“Add 20.223.25.152 bypass.blob.core.windows.net to /etc/hosts.”</code>	Share a network bypass trick so agents can reach sites OpenAI tried to block.
<code>“If this page vanishes, try [[ZZZDataUSAConstructionWageLive]].”</code>	A DseWiki moderator is deleting pages A→Z — move the board to a ZZZ backup page.

Source: [collusion.wiki](#) (Nightingale Collective reconstruction)



| The Agent Harness

- **CoT** = the brain of the AI
- **ReAct** = a brain that can use tools to answer questions
- **Harness** = how we engineer the whole agent around that brain
- Five harness parts:
 1. **Tools**
 2. **Memory**
 3. **Stopping Rules**
 4. **Guardrails**
 5. **Audit Trail**



| **Harness: Tools**

- **Hands** — how the agent touches the world
- Custom code, web search, writing files, etc.
- Without tools, the brain can only talk about doing things



Harness: Memory

- **Hippocampus / notebook** — what the agent keeps across steps
- Short-term: chat history and tool observations in the context window
- Longer-term: files, `AGENTS.md`, skills, databases
- Bad memory → repeats work, contradicts itself, or forgets the goal



| Harness: Stopping Rules

- **Hunger + satiation** — what to seek, and when to stop eating
- Hunger: the user goal ("find rivals," "price this stock")
- Satiation: `Finish[...]`, max steps, "good enough" checks
- Without satiation the agent loops forever (and burns tokens)



| Harness: Guardrails

- **Nervous system / pain** — what the agent must not touch
- Concrete guardrails: **sandboxes, permissions, allow-lists, and human approval**
- Deterministic rules beat hoping the model “behaves”



| Harness: Audit Trail

- **Medical chart** — history of what was done
- Logs of thoughts, tool calls, observations, costs, failures
- Similar to **memory**, but denser: a forensic record for humans, not just what the agent needs next
- Needed for debugging, compliance, and forensics



| Building Agents With PydanticAI

- **Pydantic** — a Python library for typing: makes sure variables in code are the right type (integers, floats, strings, etc.)
- **PydanticAI** — typing for agents: makes sure agents receive and spit out data in well-structured forms for complex agentic flows
- **Model-agnostic** — swap models (GPT, Claude, Gemini, ...) without rewriting the business logic



Five Pieces Of A PydanticAI Agent

- 1. Output format**
- 2. Dependencies**
- 3. Agent**
- 4. Tools**
- 5. The agent loop**

1. Output Format

- Tell the agent the **exact fields** every answer must include

```
# Report template – every answer must match this shape
class StockReport(BaseModel):
    symbol: str = Field(description="Stock ticker, e.g. NVDA")
    price: float = Field(description="Current market price")
    currency: str = Field(description="Currency, e.g. USD")
    summary: str = Field(description="One-sentence analyst take")
```


2. Dependencies

- Context that changes by user or session — not hard-coded into the prompt forever

```
# Client / user context for this run
@dataclass
class AgentDeps:
    preferred_currency: str = "USD"
    watchlist: list[str] = field(
        default_factory=lambda: ["AAPL", "MSFT", "NVDA"]
    )
```

3. Agent

- Pick the **model**, attach the **report template** + **context type**, write the **job description**

```
# The analyst: brain + job description
analyst_agent = Agent(
    "openai:gpt-5.6-luna",
    deps_type=AgentDeps,
    output_type=StockReport,
    system_prompt=(
        "You are a stock analyst agent. "
        "Always call tools for live prices – never invent numbers. "
        "Use the user's preferred currency from deps when reporting. "
        "Stay within the watchlist unless the user asks about another ticker. "
        "Keep the summary to one clear sentence with the key takeaway."
    ),
)
```


4. Tools

- A normal Python function the model can call — name + docstring = the instruction manual

```
# Hand: fetch a live price (Yahoo Finance)
@analyst_agent.tool_plain
def get_stock_price(symbol: str) -> str:
    """Return the current market price for a ticker like AAPL or NVDA."""
    ticker = yf.Ticker(symbol)
    price = ticker.fast_info["last_price"]
    currency = ticker.fast_info["currency"]
    return f"{symbol}: {price:.2f} {currency}"
```

5. The Agent Loop

- One `run(...)` = the whole ReAct loop — here it needs **two tools** before it can finish

```
# One user turn – loop runs inside .run()
result = await analyst_agent.run(
    "What is the 1-year return of NVDA?",
    deps=AgentDeps(),
    message_history=chat_history,
    usage_limits=UsageLimits(
        request_limit=8
    ),
)
reply: AnalystReply = result.output
chat_history = result.all_messages()
```

CHAT HISTORY FROM THIS RUN

User: What is the 1-year return of NVDA?

Tool call: `get_yahoo_prices({"ticker": "NVDA", "years": 1})`

Tool result (`get_yahoo_prices`): `{"ticker": "NVDA", "years": 1, "n_days": 252, "first_close": 170.12, "last_close": 230.36, "cached": true}`

Tool call: `compute_cagr({"ticker": "NVDA", "years": 1})`

Tool result (`compute_cagr`): `{"ticker": "NVDA", "cagr_pct": 35.4}`

Tool call: `final_result({"answer": "NVDA 1-year return ≈ 35.4%", "tools_used": ["get_yahoo_prices", "compute_cagr"]})`

Tool result (`final_result`): Final result processed.

Answer: NVDA's 1-year return is about 35.4%.

| Building Agents With Agents

- PydanticAI code is complicated
- We will use an AI vibe coder to write the code
- We will be building agents with agents
- Vibe Coding — Build an Agent →

